GPU comparison chart

		GPU RENDER	GPU L4	GPU L40S	GPU H100 PCIe Single (and dual)	Jero (2DGX H100, 16 H100 GPUs)	Nabu (127 DGX H100, 1016 H100 GPUs)
Type		Instance	Instance	Instance	Instance	Supercomputer	Supercomputer
Training	NLP, LLM	+	X	++	+++	++++	++++
	GenAI (Image, video)	X	X	++	+++	++++	++++
	Recommender System	+	X	++	+++	++++	++++
Inference	NLP, LLM		+ (up to 5b parameters)	++ (up to 65b parameters) and + from 66b to 175b	+++	++++	X (training cluster not made for inference)
	Generative AI (image, video)		+	+++	++	X (training cluster not made for inference)	X (training cluster not made for inference)
	Graph, vector Database		X	X	X (not the best price/ performance)	X (training cluster not made for inference)	X (training cluster not made for inference)
	Recommender System		X	X	+++	X (training cluster not made for inference)	X (training cluster not made for inference)
	AI Video		+++	++	X (not made for this)	X (training cluster not made for inference)	X (training cluster not made for inference)
	Computer vision		+++	++	X (not the best price/ performance)	X (training cluster not made for inference)	X (training cluster not made for inference)
HPC		X	X	X	X	++	++++
Rendering		+ (up to 4K)	+	++	X	X	X
Remote Desktop (full GPU)		+	+++	++	X	X	X
Virtual Desktop (VDI)			+++	X	X	X	X
AI Video			+++	++	X	X	X